



源來適你

OpenSource4You

since 2022

From Taiwan to the World — Open Source Community



Who Are We?

A Loose Organization, Serious About Open Source



OpenSource4You is a loosely organized but deeply passionate open source community. We bring together engineers, students, and industry professionals who care about open source. Our goal is to "let the community empower the individual" — through real contributions to open source projects, we help every member shine on the global stage.

large-project
20⁺
maintainers cultivated

every month
40⁺
online events

every month
20,000⁺
technical discussions



every year
30 +
in-person events

What We Do



Mentorship Program

Senior contributors guide new members one-on-one, helping them contribute to real international open source projects from the ground up. No matter your current level, you can progress step by step toward becoming a core open source developer.



Technical Activities

- "Tech Talks" — online technical sharing sessions
- Study groups on Kafka, Ray, Flyte, and more
- In-person meetups



Content Production


- Facebook community posts
- GitHub technical notes
- "OpenSource4You" podcast



International Engagement

- Technical talks at international conferences
- Taiwan's first Apache Software Foundation chapter





Apache Kafka

The Undefeated Legend of the Modern Data World

Kafka is the absolute standard for distributed event streaming worldwide. Over 80% of Fortune 100 companies rely on it to run their core infrastructure; Netflix, Uber, and LinkedIn process trillions of real-time messages with it daily — rock-solid and never down.

- Million-messages-per-second throughput with industry-leading low latency
- Highly reliable architecture with zero data loss
- KRaft architecture upgrade makes cluster management easier

**When the world's engineers face extreme concurrency,
this is their one shared faith.**



Apache Airflow

The Air Traffic Control Tower for Data Pipelines

Apache Airflow is one of the most widely used workflow scheduling and monitoring platforms in the world. Engineers define pipelines in Python as a first-class citizen, scheduling, monitoring, and managing complex cross-system workflows with ease. From daily reporting to ML training, Airflow handles it end-to-end.

- Python-first workflow definitions, minimal learning curve for engineers
- Rich operator ecosystem, integrating hundreds of databases and cloud services
- AI-native integration with smart, agent-in-the-loop retries.
- Full-featured web UI with clear visibility into pipeline status.

**An essential weapon for data engineers
— the industry standard, no question.**



RAY

The Compute Brain of the Big AI Era

Ray is the next-generation operating system for AI infrastructure. OpenAI uses it to orchestrate tens of thousands of GPUs to train GPT-4; with just one Python decorator, `@ray.remote`, you can scale a program from your laptop into a distributed system commanding thousands of servers. NVIDIA, AWS, and Google Cloud all race to integrate it natively — it sits at the core of the entire AI ecosystem.

- Distributed scaling with a single command — developer-friendly
- The industry's first choice for large-scale model training and inference
- Tech giants compete to integrate it as the ecosystem keeps growing

The compute faith of every AI engineer.



Apache Ozone

The Infinite Warehouse of the Cloud-Native Era

Ozone is a distributed object storage system built for massive-scale deployments — the next-generation storage core of the Hadoop ecosystem. It breaks through the bottleneck traditional HDFS faces with huge numbers of small files, scaling effortlessly to billions of objects and hundreds of petabytes in cloud-native Kubernetes environments, all while maintaining extremely high throughput and availability.

- Supports billions of objects, fully solving small-file performance issues
- Kubernetes-native, with seamless hybrid cloud and on-premises deployment
- S3 API compatible — existing tools plug in with zero changes

Letting enterprise data lakes truly become boundless and bottomless.



Flyte

The Workflow Engine That Precisely Commands Billions of Tasks

Open-sourced by Lyft, Flyte is a Kubernetes-native workflow engine purpose-built for machine learning and data pipelines. Top companies like Spotify use it to precisely orchestrate complex AI workflows. Its intelligent caching mechanism lets tasks resume directly from the point of interruption, saving substantial recomputation costs.

- Strongly typed design with complete, traceable data lineage
- GitOps integration — pipeline versions are controllable with one-click rollback
- Kubernetes-native architecture, born for the cloud

Freeing ML engineers from all-night overtime.



Apache DataFusion

The Query-Speed Legend Built in Rust

DataFusion is a high-performance embeddable query engine built in Rust — a core component of the Apache Arrow ecosystem. Its vectorized execution engine squeezes every ounce of power from modern CPUs, making SQL queries astonishingly fast. Next-generation data platforms like InfluxDB IOx and Delta Lake use DataFusion as their underlying query brain.

- Native Rust implementation, far outperforming traditional JVM query engines
- Vectorized execution paired with the Arrow memory format for ultimate IO efficiency
- Embeddable into any system — the best foundation for building custom query engines

**The shared foundation of next-generation data tools,
quietly rewriting the industry's rules.**



Apache YuniKorn

The Intelligent Scheduling Brain of Kubernetes Clusters

YuniKorn is a next-generation Kubernetes resource scheduler designed for large-scale AI/ML workloads and multi-tenant enterprise environments. It replaces the limitations of the default scheduler with hierarchical queue management, maximizing cluster resource utilization and eliminating idle compute waste entirely.

- Gang Scheduling support ensures distributed training tasks launch as a complete batch
- Hierarchical multi-tenant queues fairly allocate cluster resources — no more scrambling
- Seamless integration with native Kubernetes APIs — zero changes to existing workflows

Letting every GPU and CPU do the right thing at the right time.



Apache Mahout

Building High-Performance Environments for Quantum Machine Learning

Mahout provides a cutting-edge environment for building scalable, performant machine learning applications, with its latest evolution focusing entirely on the frontier of quantum computing. Powered by a high-performance Rust and CUDA backend, Mahout seamlessly bridges classical data science and quantum workflows, allowing developers to write quantum applications once and execute them anywhere.

- Providing unified Python API to run circuits across Qiskit, Cirq, and Amazon Braket.
- Ultra-fast, GPU-accelerated encoding of classical data into quantum states.
- Built-in DLPack support enabling zero-copy tensor transfers.

**When machine learning meets quantum computing,
Mahout is writing the next chapter.**

The Cornerstones of the Open Source World



Chia-Ping Tsai
Apache Member



Wei-Chiu Chuang
Apache Member



PoAn Yang
Apache Kafka
Committer



Kai-Hsun Chen
RAY
Maintainer



Kuan-Po Tseng
Apache YuniKorn
Committer



Tingyao Huang
Apache YuniKorn
PMC Member



Chung-En Lee
Apache Ozone
PMC Member



Wei Lee
Apache Airflow
PMC Member



Kevin Su
Flyte
Maintainer



TengYao Chi
Apache Kafka
Committer



Eric Chang
Apache Gravitino
Committer



Zhe-You Liu
Apache Airflow
Committer




Jie-Kai Chang

Apache Mahout
PMC Member



Chu Cheng Li

Apache Ozone
Committer



Jheng-Sing Chen

Apache Ozone
Committer



Kuan-Hao Huang

Apache Mahout
Committer



Hsien-Cheng Huang

Apache Mahout
Committer



Guan-Ming Chiu

Apache Mahout
PMC Member



Han-Ju Chen

RAY
Maintainer



Nary Yeh

Flyte
Maintainer



Jiunn-Yang Huang

Apache Kafka
Contributor



Yu-Lin Chen

Apache YuniKorn
Committer



Huang-Hsiang Cheng

Apache DataFusion
Comet Contributor



Jui-An Huang

RAY
Maintainer



Alex Wu

Flyte
Maintainer



Yu-Chen Lai

Apache DataFusion
Comet Contributor




Ming-Yen Chung

Apache Kafka
Contributor



You-Cheng Lin

RAY
Maintainer



Taiwan Open Source Starts Here

